

CS 378 (Spring 2003)

Linux Kernel Programming

Yongguang Zhang

(ygz@cs.utexas.edu)

This Lecture

- Time Management
- Synchronization
- Questions?

Clocks in Linux

- Real-Time Clock (RTC)
 - Independent of CPU
 - May be expensive to access
- CPU Cycle Counter (TSC Register)
 - For accurate time measurements
 - Granularity: CPU cycle (e.g. 1GHz)
- Kernel Clock (Programmable Interval Timer)
 - Each tick: cause timer interrupt at a fixed interval
 - For the kernel to keep track of time
 - Granularity: much larger: in terms of 10 Hz-1000 Hz

CPU Cycle Counter

- Architecture dependent
 - i386: TSC (Time Stamp Counter) register, 64 bits
 - Updated by hardware, read by "rdtsc" instruction
- C functions
 - In `include/asm-i386/{msr.h,timex.h}`
`rdtsc()`, `rdtscl()`, `rdtscll()`, `cycles_t get_cycles (void)`
- Example uses:
 - `Gettimeofday()`
 - Networking (timestamping, scheduling)
 - Some device drivers

Kernel Timer

- Frequency: HZ (per second)
 - Defined in include/asm/param.h
 - Linux chooses 100Hz: #define HZ 100
- **jiffies**: number of ticks since system boot time
 - Global variable declared in include/linux/sched.h
extern unsigned long volatile jiffies;
 - Incremented in each timer interrupt
 - 32-bit, so **jiffies** can wrap around
 - Never: if (jiffies > last_jiffies)
 - Always: if ((long)(jiffies – last_jiffies) > 0)

Timer Interrupt Handler

- `timer_interrupt()`
 - Architecture-dependent: `arch/i386/kernel/time.c`
 - Save TSC value (lower 32bit) into `last_tsc_low`
 - Then calls `do_timer_interrupt()`
- `do_timer_interrupt()`
 - Acks the interrupt (if necessary)
 - Calls the architecture independent `do_timer()`
 - Updates the RTC every 660 ticks if configured

do_timer()

- Architecture independent timer interrupt routine
 - Increments jiffies
 - Update process time (add 1 tick to user or system time)
 - Schedule timer bottom half
 - Essentially, (in kernel/timer.c)

```
(*(unsigned long *)&jiffies)++;  
update_process_times(user_mode(regs));  
mark_bh(TIMER_BH);
```
 - Q: why is jiffies updated in “top half”?

Timer Button Half

- A high priority tasklet to run at a later time
 - Button Half: backward compatible only, will go away
 - `TIMER_BH` → `timer_bh()`
- `timer_bh()`
 - Architecture independent, in `kernel/timer.c`
 - Calls the following two functions
 - `update_times()`: update the coarse wall clock (`xtime`)
 - `run_timer_list()`: process the list of kernel event timers

Wall Clock Management

- Wall clock
 - Time-of-day clock, real-time clock
 - To read precise wall clock: `do_gettimeofday()`
 - Two parts (each a long integer)
 - Second since 1/1/1970 (valid until 1/19/2038 03:14:07 UTC)
 - Microsecond within the second
- Implementation
 - Read hardware RTC only when system boot
 - Kernel maintains a coarse clock (`xtime`)
 - Accurate value calculated with help from TSC

xtime

- Coarse timer, updated in timer BH only
 - Defined in include/linux/sched.h
 - Two fields: xtime.tv_sec, xtime.tv_usec
 - timer_bh() calls update_times(), which does

```
ticks = jiffies - wall_jiffies;
if (ticks) {
    wall_jiffies += ticks;
    update_wall_time(ticks);
}
```
 - Q1: what is the actual granuty of xtime.tv_usec?
 - Q2: can xtime drift?

do_gettimeofday()

- Return precise time of day clock value (usec)
 - Need to compensate `xtime.tv_usec`
 - Add the lost jiffies: `jiffies - wall_jiffies`;
 - Add the elapsed cycles: read TSC value and substrate `last_tsc_low`
- Use of `do_gettimeofday()`
 - Serving system call `gettimeofday()`
 - Other places need accurate time of day (e.g. logging)
- Kernel programming
 - Use `do_gettimeofday()` or `get_cycles()`
 - Don't use `xtime` directly

Kernel Event Timers

- Function calls scheduled at a specified time
- Data structure
 - In kernel/linux/timer.h

```
struct timer_list {  
    struct list_head list;  
    unsigned long expires;  
    unsigned long data;  
    void (*function)(unsigned long);  
};
```
 - expires: function call at this jiffies value
 - function, data: function and argument to call

Using Kernel Timers

- Create a `struct timer_list` object
 - Add initialize it with `init_timer(struct timer_list *)`
- Setup the fields
- Insert it into the kernel timer list: `add_timer()`
- Before expiration, you can
 - Reschedule it: `mod_timer(t, new_expires)`
 - Delete timer: `del_timer_sync()`

Event Timer Handling

- Checking and executing event timers
 - By `run_timer_list()`
 - During timer BH only, so timing can be imprecise
- Kernel data structure for managing event timers
 - Group events into 512 lists by their expiration time
 - First 256 lists: events expiring in next 1, 2, 3, ..., 256 ticks
 - Next 64 lists: in next 1×2^8 , 2×2^8 , 3×2^8 , ..., 64×2^8 ticks
 - Next 64 lists: in next 1×2^{14} , 2×2^{14} , 3×2^{14} , ..., 64×2^{14} ticks
 - Next 64 lists: in next 1×2^{20} , 2×2^{20} , 3×2^{20} , ..., 64×2^{20} ticks
 - Last 64 lists: in next 1×2^{26} , 2×2^{26} , 3×2^{26} , ..., 64×2^{26} ($=2^{32}$) ticks
 - Code in `kernel/timer.c`

Deferred Execution

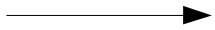
- Scheduling execution in a safe time, ASAP
 - Use tasklet
 - In interrupt context (task cannot sleep)
 - Mostly used by interrupt handler/device driver
- Scheduling execution in a specific time
 - Use kernel event timer
 - Executed by timer BH, which is a tasklet afterall – still in interrupt context, cannot sleep
- Create your own list and scheduling
 - Use task queue (may go away in 2.6 kernel)

Task Queues

- Use task queue (see `include/linux/tqueue.h`)
 - Declare a queue
`DECLARE_TASK_QUEUE(name)`
 - Add task to a queue
`int queue_task(struct tq_struct *, task_queue *)`
 - Run all tasks currently in the queue
`void run_task_queue(task_queue *)`
- Task queue may run in a different context
 - May run in interrupt context (by tasklet)
 - May run in a different thread (e.g., `ksoftirqd`, `keventd`)
 - May run in a different CPU

Time Illustration

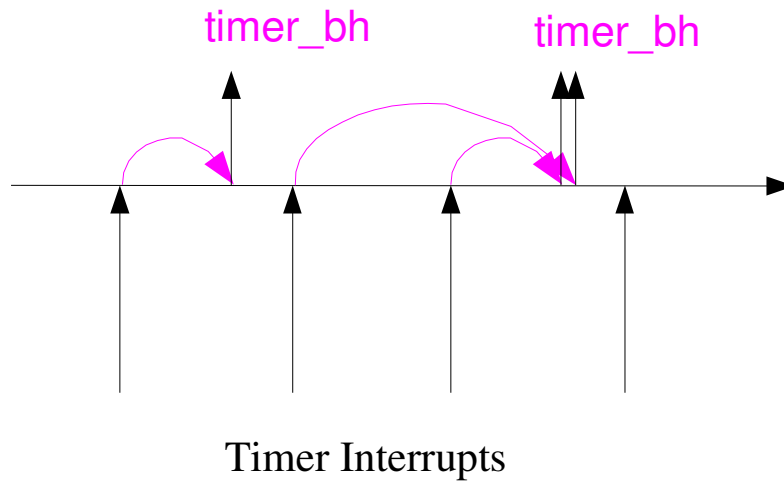
User mode



Kernel mode



Interrupt mode



Synchronization

- You all know why we need synchronization
 - Whenever a resource can be accessed by others
 - This task can sleep and schedule to others
 - This task can be interrupted and tasklet can run
 - There are other tasks in other CPUs
- Synchronization mechanisms in Linux kernel
 - Atomic operations
 - Disabling interrupt
 - Locking: spin locks, semaphores

Hardware Support

- Fundamentally, mutual exclusion requires hardware support
 - Bootstrap from hardware-supported atomic action
- Single processor
 - cli and sti instruction: disable (and enable) all interrupts
- SMP architecture
 - The “lock” instruction prefix: lock the memory bus for this instruction (so no other CPU can access the memory until this instruction is done)

Atomic Operations

- Execute a single instruction in an “atomic” way, even under multiprocessor system
 - Supported by SMP hardware (with the “lock” instruction prefix)
 - Two types: bit ops and atomic integer variables
- Bit ops: change a bit in any memory address
 - `void set_bit(int nr, volatile void * addr)`
 - `void clean_bit(int nr, volatile void * addr)`
 - `int test_and_set_bit(int nr, volatile void * addr)`
 - And more (see `include/asm-i386/bitops.h`)

Atomic Integer Variable

- Atomic operations on integers
 - Defined in `include/asm-i386/atomic.h`
 - Data Structure: type `atomic_t`
- Functions:
 - `atomic_read(v)`
 - `atomic_set(v,i)`
 - `atomic_add(i,v)`
 - `atomic_inc(v)`
 - ...

Interrupt Disabling

- Disable/enable all interrupts in this CPU
 - Use `local_irq_disable()` and `local_irq_enable()`
 - Implemented by `cli` or `sti` instruction
 - In SMP system: has no effect on other CPUs
- Global interrupt disabling/enabling
 - Use `cli()` and `sti()`
 - In uniprocessor system: same as `cli` and `sti` instruction
 - In SMP, `cli()` and `sti()` are implemented with spin lock to delay interrupt handlers in other CPUs

Saving eflags Register Content

- Must save register content before interrupt disabling and restore it after re-enabling
 - Register includes the interrupt flag (IF)
 - See `include/asm-i386/system.h`
- For local interrupt disabling
 - `__save_flags(long)` and `__restore_flags(long)`
- For global interrupt disabling
 - `save_flags(long)` and `restore_flags(long)`

Protecting Critical Session

- Using interrupt disabling (the simple way)

...

unsigned long flags;

save_flags(flags);

cli();

... critical session ...

restore_flags(flags);

...

- Disadvantage: stops all other CPUs,
- Protecting from deferred tasklet/BH only?
 - Use `local_bh_disable()`, `local_bh_enable()`

Spin Lock

- A locking mechanism for SMP system
 - Through a shared variable
 - Acquire the lock by setting the variable
 - “Spin” in a busy-wait loop until the variable is unset
 - Should be used with care: holding a spin lock too long may cause other CPUs to waste time in busy waiting
 - Data type: `spinlock_t` (in `include/asm/spinlock.h`)
- In a uniprocessor system
 - Implemented as no-op (because it is the only process running)

Using Spin Lock

- Include `<linux/spinlock.h>`
- Define spin lock variable:
 - `spinlock_t my_lock = SPIN_LOCK_UNLOCKED;`
- To lock, call `spin_lock(my_lock)`
 - With local interrupt disabled: `spin_lock_irq()`
 - Also with eflags also: `spin_lock_irqsave()`
 - With tasklet/BH disabled: `spin_lock_bh()`
- To unlock, call `spin_unlock(my_lock)`
- To check, `spin_is_locked(my_lock)` returns 1/0

– `void spin_trylock()` `spin_unlock_wait()`

Read/Write Spin Lock

- Allow multiple readers but only one writer
- Data type: `rwlock_t`
 - In `include/spinlock.h`
 - Ex: `rwlock_t my_lock = RW_LOCK_UNLOCKED;`
- Operations:
 - `void read_lock(rwlock_t *rw)`
 - `void read_unlock(rwlock_t *rw)`
 - `void write_lock(rwlock_t *rw)`
 - `void write_unlock(rwlock_t *rw)`
 - And more (with `_irq`, `_irqsave`, `_bh`, ...)

Kernel Semaphores

- Concept of Semaphore
 - A number of resource available for claimed by task
 - Task put on the wait queue if resource unavailable
 - Task waits up when resource available (released)
- Data Structure
 - In include/asm/semaphore.h

```
struct semaphore {  
    atomic_t count;    // > 0: available, <=0: busy  
    int sleepers;  
    wait_queue_head_t wait;  
};
```

Using Kernel Semaphores

- To Initialize: `sema_init(struct semaphore *, int)`
 - If number of resource is 1, it is called MUTEX
 - `init_MUTEX(sem)`, `init_MUTEX_LOCKED(sem)`
- To use a resource: `down(struct semaphore * sem)`
 - Atomically decrease count
 - Put current on the wait queue if `count < 0`
 - Variants: `down_trylock()`, `down_interruptible()`
- To release: `up(struct semaphore * sem)`
 - Atomically increase count
 - Wake up one task on wait queue if `count <= 0`

Synchronization Mechanisms

- Atomic operations
- Disabling interrupts
 - Simple way, but reduce parallelism
- Spin locks
 - Better way, good for protecting against other CPUs
- Semaphores
 - Good for synchronization among different tasks
 - Invoke scheduling, cannot be used in interrupt context

Summary

- Time Management:
 - LKP §3.2.5
 - ULK §6
 - LDD2 §7
- Synchronization:
 - LKP §5.1, §10
 - ULK §5
 - LDD2 §9
- Next Lecture: Device Driver